

ASSIGNMENT 2

Kulsoom Khurshid

SP20-BCS-044

TASK 1

This dataset is related to flowers. There are four characteristics i.e., sepal length, sepal width, petal length and petal width. On the basis of these characteristics classification is made whether the flower is Setosa, Versicolor and Virginica. This dataset is divided in two, half data is used to train the model whereas the other half is used for testing to check if the flower is correctly categorized.

import seaborn as sns

This library that I have used to load the Iris dataset.

The two algorithms that we have used are Logistic regression and Decision Tree.

```
In [28]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
import os
os.getcwd()
```

```
Out[28]: 'D:\\Machine Learning'
```

```
In [8]: # We are reading our data
df=sns.load_dataset("iris")
df.head()
```

```
Out[8]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa

TASK 2

df.describe()

```
In [2]: # We are reading our data
df=sns.load_dataset("iris")
df.describe()
```

```
Out[2]:
```

	sepal_length	sepal_width	petal_length	petal_width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333
std	0.828066	0.435866	1.765298	0.762238
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

df.info()

```
In [3]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   sepal_length    150 non-null   float64
1   sepal_width     150 non-null   float64
2   petal_length    150 non-null   float64
3   petal_width     150 non-null   float64
4   species         150 non-null   object  
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

TASK 3

Logistic Regression

```
In [10]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
import seaborn as sns
import os
os.getcwd()
```

Out[10]: 'D:\Machine Learning'

```
In [11]: # We are reading our data
df=sns.load_dataset("iris")
df.head()
```

Out[11]:

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa

```
In [12]: from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(df.iloc[:, :-1], df.iloc[:, -1], test_size=0.2, random_state=2)
```

```
In [13]: from sklearn.linear_model import LogisticRegression
```

```
In [14]: clf1=LogisticRegression()
```

```
In [15]: clf1.fit(X_train, Y_train)
```

Out[15]: LogisticRegression()

```
In [16]: y_pred1 = clf1.predict(X_test)
```

```
In [19]: from sklearn.metrics import accuracy_score, confusion_matrix
print("Accuracy of Logistic Regression", accuracy_score(Y_test, y_pred1))
```

Accuracy of Logistic Regression 0.9666666666666667

```
In [20]: confusion_matrix(Y_test, y_pred1)
```

Out[20]: array([[14, 0, 0],
[0, 7, 1],
[0, 0, 8]], dtype=int64)

```
In [21]: result = pd.DataFrame()
result['Actual Label'] = Y_test
result['Logistic Regression Prediction'] = y_pred1
```

```
In [22]: result.sample(10)
```

```
Out[22]:
```

	Actual Label	Logistic Regression Prediction
74	versicolor	versicolor
115	virginica	virginica
113	virginica	virginica
29	setosa	setosa
87	versicolor	versicolor
127	virginica	virginica
6	setosa	setosa
77	versicolor	virginica
25	setosa	setosa
128	virginica	virginica

```
In [23]: from sklearn.metrics import recall_score,precision_score,f1_score
```

```
In [24]: print("For Logistic regression Model")
print("-"*50)
cdf = pd.DataFrame(confusion_matrix(Y_test,y_pred1),columns=list(range(0,3)))
print(cdf)
print("-"*50)
print("Precision - ",precision_score(Y_test,y_pred1,average='macro'))
print("Recall - ",recall_score(Y_test,y_pred1,average='macro'))
print("F1 score - ",f1_score(Y_test,y_pred1,average='macro'))
```

```
For Logistic regression Model
-----
      0   1   2
0  14   0   0
1   0   7   1
2   0   0   8
-----
Precision -  0.9629629629629629
Recall -    0.9583333333333334
F1 score -  0.9581699346405229
```

```
In [25]: precision_score(Y_test,y_pred1,average='macro')
```

```
Out[25]: 0.9629629629629629
```

```
In [26]: from sklearn.metrics import classification_report
print (classification_report(Y_test, y_pred1))
```

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	14
versicolor	1.00	0.88	0.93	8
virginica	0.89	1.00	0.94	8
accuracy			0.97	30
macro avg	0.96	0.96	0.96	30
weighted avg	0.97	0.97	0.97	30

```
In [28]: plt.xlabel('Features')
plt.ylabel('Species')


pltX = df.loc[:, 'sepal_length']
pltY = df.loc[:, 'species']
plt.scatter(pltX, pltY, color='red', label='sepal_length')

pltX = df.loc[:, 'sepal_width']
pltY = df.loc[:, 'species']
plt.scatter(pltX, pltY, color='blue', label='sepal_width')

pltX = df.loc[:, 'petal_length']
pltY = df.loc[:, 'species']
plt.scatter(pltX, pltY, color='yellow', label='petal_length')

pltX = df.loc[:, 'petal_width']
pltY = df.loc[:, 'species']
plt.scatter(pltX, pltY, color='green', label='petal_width')

plt.legend(loc=4, prop={'size':8})
plt.show()
```



TASK 4

Decision Tree

```
In [18]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
import os
os.getcwd()
from sklearn.datasets import make_classification
```

```
In [19]: x,y= make_classification(n_features=4, n_redundant=0, n_informative=4, n_clusters_per_class=1)
```

```
In [20]: # We are reading our data
df=sns.load_dataset("iris")
df.head()
```

```
Out[20]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa

```
In [21]: # function for row sampling
def sample_rows(df, percent):
    return df.sample(int(percent*df.shape[0]))
```

```
In [22]: #function for column sampling
def sample_features(df, percent):
    cols= random.sample(df.columns.tolist()[:-1], int(percent*df.shape[1]))
    return df[cols]
```

```
In [23]: #function for combined sampling
def combined_sampling(df, row_percent, col_percent):
    new_df=sample_rows(df, row_percent)
    return sample_features(new_df, col_percent)
```

```

In [24]: df1=sample_rows(df, 0.1)

In [25]: df2=sample_rows(df, 0.1)

In [26]: df3= sample_rows(df, 0.1)

In [27]: df3.shape
Out[27]: (15, 5)

In [28]: from sklearn.tree import DecisionTreeClassifier
         clf1= DecisionTreeClassifier()
         clf2= DecisionTreeClassifier()
         clf3= DecisionTreeClassifier()

In [32]: clf1.fit(df1.iloc[:,0:4],df1.iloc[:,-1])
         clf2.fit(df2.iloc[:,0:4],df1.iloc[:,-1])
         clf3.fit(df3.iloc[:,0:4],df1.iloc[:,-1])

Out[32]: DecisionTreeClassifier()

```

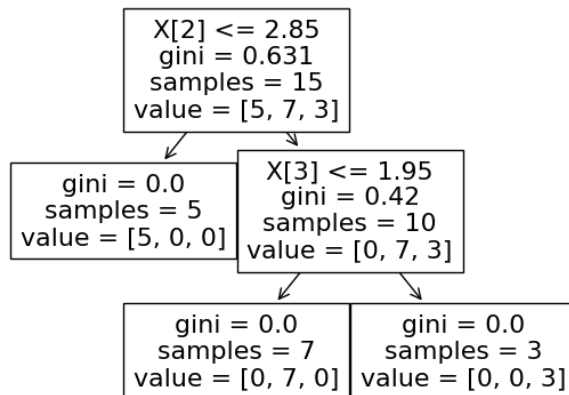
```

In [33]: from sklearn.tree import plot_tree

In [34]: plot_tree(clf1)

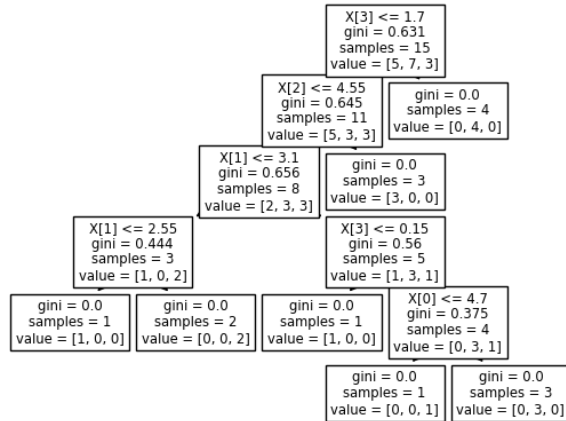
Out[34]: [Text(0.4, 0.8333333333333334, 'X[2] <= 2.85\ngini = 0.631\nsamples = 15\nvalue = [5, 7, 3]'),
Text(0.2, 0.5, 'gini = 0.0\nsamples = 5\nvalue = [5, 0, 0]'),
Text(0.6, 0.5, 'X[3] <= 1.95\ngini = 0.42\nsamples = 10\nvalue = [0, 7, 3]'),
Text(0.4, 0.16666666666666666, 'gini = 0.0\nsamples = 7\nvalue = [0, 7, 0]'),
Text(0.8, 0.16666666666666666, 'gini = 0.0\nsamples = 3\nvalue = [0, 0, 3]')]

```



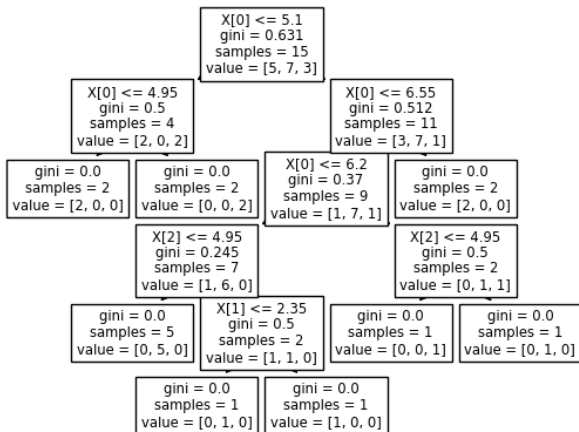
In [35]: plot_tree(clf2)

```
Out[35]: [Text(0.6666666666666666, 0.9166666666666666, 'X[3] <= 1.7\ngini = 0.631\nsamples = 15\nvalue = [5, 7, 3]'),
Text(0.5555555555555556, 0.75, 'X[2] <= 4.55\ngini = 0.645\nsamples = 11\nvalue = [5, 3, 3]'),
Text(0.4444444444444444, 0.5833333333333334, 'X[1] <= 3.1\ngini = 0.656\nsamples = 8\nvalue = [2, 3, 3]'),
Text(0.2222222222222222, 0.4166666666666667, 'X[1] <= 2.55\ngini = 0.444\nsamples = 3\nvalue = [1, 0, 2]'),
Text(0.1111111111111111, 0.25, 'gini = 0.0\nsamples = 1\nvalue = [1, 0, 0]'),
Text(0.3333333333333333, 0.25, 'gini = 0.0\nsamples = 2\nvalue = [0, 0, 2]'),
Text(0.6666666666666666, 0.4166666666666667, 'X[3] <= 0.15\ngini = 0.56\nsamples = 5\nvalue = [1, 3, 1]'),
Text(0.5555555555555556, 0.25, 'gini = 0.0\nsamples = 1\nvalue = [1, 0, 0]'),
Text(0.7777777777777778, 0.25, 'X[0] <= 4.7\ngini = 0.375\nsamples = 4\nvalue = [0, 3, 1]'),
Text(0.6666666666666666, 0.0833333333333333, 'gini = 0.0\nsamples = 1\nvalue = [0, 0, 1]'),
Text(0.8888888888888888, 0.0833333333333333, 'gini = 0.0\nsamples = 3\nvalue = [0, 3, 0]'),
Text(0.6666666666666666, 0.5833333333333334, 'gini = 0.0\nsamples = 3\nvalue = [3, 0, 0]'),
Text(0.7777777777777778, 0.75, 'gini = 0.0\nsamples = 4\nvalue = [0, 4, 0]')]
```



In [36]: plot_tree(clf3)

```
Out[36]: [Text(0.4444444444444444, 0.9166666666666666, 'X[0] <= 5.1\ngini = 0.631\nsamples = 15\nvalue = [5, 7, 3]'),
Text(0.2222222222222222, 0.75, 'X[0] <= 4.95\ngini = 0.5\nsamples = 4\nvalue = [2, 0, 2]'),
Text(0.1111111111111111, 0.5833333333333334, 'gini = 0.0\nsamples = 2\nvalue = [2, 0, 0]'),
Text(0.3333333333333333, 0.5833333333333334, 'gini = 0.0\nsamples = 2\nvalue = [0, 0, 2]'),
Text(0.6666666666666666, 0.75, 'X[0] <= 6.55\ngini = 0.512\nsamples = 11\nvalue = [3, 7, 1]'),
Text(0.5555555555555556, 0.5833333333333334, 'X[0] <= 6.2\ngini = 0.37\nsamples = 9\nvalue = [1, 7, 1]'),
Text(0.3333333333333333, 0.4166666666666667, 'X[2] <= 4.95\ngini = 0.245\nsamples = 7\nvalue = [1, 6, 0]'),
Text(0.2222222222222222, 0.25, 'gini = 0.0\nsamples = 5\nvalue = [0, 5, 0]'),
Text(0.4444444444444444, 0.25, 'X[1] <= 2.35\ngini = 0.5\nsamples = 2\nvalue = [1, 1, 0]'),
Text(0.3333333333333333, 0.0833333333333333, 'gini = 0.0\nsamples = 1\nvalue = [0, 1, 0]'),
Text(0.5555555555555556, 0.0833333333333333, 'gini = 0.0\nsamples = 1\nvalue = [1, 0, 0]'),
Text(0.7777777777777778, 0.4166666666666667, 'X[2] <= 4.95\ngini = 0.5\nsamples = 2\nvalue = [0, 1, 1]'),
Text(0.6666666666666666, 0.25, 'gini = 0.0\nsamples = 1\nvalue = [0, 0, 1]'),
Text(0.8888888888888888, 0.25, 'gini = 0.0\nsamples = 1\nvalue = [0, 1, 0]'),
Text(0.7777777777777778, 0.5833333333333334, 'gini = 0.0\nsamples = 2\nvalue = [2, 0, 0]')]
```



TASK 5

```
In [78]: from sklearn.metrics import accuracy_score, confusion_matrix
print("Accuracy of Logistic Regression", accuracy_score(Y_test, y_pred1))
print("Accuracy of Decision Trees", accuracy_score(Y_test, y_pred2))
```

```
Accuracy of Logistic Regression 0.9666666666666667
Accuracy of Decision Trees 0.9333333333333333
```

```
In [80]: print("Logistic Regression Confusion Matrix\n")
pd.DataFrame(confusion_matrix(Y_test, y_pred1), columns=list(range(0,3)))
```

Logistic Regression Confusion Matrix

Out[80]:

	0	1	2
0	14	0	0
1	0	7	1
2	0	0	8

```
In [81]: print("Decision Tree Confusion Matrix\n")
pd.DataFrame(confusion_matrix(Y_test, y_pred2), columns=list(range(0,3)))
```

Decision Tree Confusion Matrix

Out[81]:

	0	1	2
0	14	0	0
1	0	7	1
2	0	1	7

```
In [85]: print("For Logistic regression Model")
print("-"*50)
cdf = pd.DataFrame(confusion_matrix(Y_test, y_pred1), columns=list(range(0,3)))
print(cdf)
print("-"*50)
print("Precision - ", precision_score(Y_test, y_pred1, average='macro'))
print("Recall - ", recall_score(Y_test, y_pred1, average='macro'))
print("F1 score - ", f1_score(Y_test, y_pred1, average='macro'))
```

For Logistic regression Model

```
-----
   0  1  2
0 14  0  0
1  0  7  1
2  0  0  8
-----
```

```
Precision -  0.9629629629629629
Recall -    0.9583333333333334
F1 score -  0.9581699346405229
```

```
In [86]: print("For DT Model")
print("-"*50)
cdf = pd.DataFrame(confusion_matrix(Y_test, y_pred2), columns=list(range(0,3)))
print(cdf)
print("-"*50)
print("Precision - ", precision_score(Y_test, y_pred2, average='macro'))
print("Recall - ", recall_score(Y_test, y_pred2, average='macro'))
print("F1 score - ", f1_score(Y_test, y_pred2, average='macro'))
```

For DT Model

```
-----
   0  1  2
0 14  0  0
1  0  7  1
2  0  1  7
-----
```

```
Precision -  0.9166666666666666
Recall -    0.9166666666666666
F1 score -  0.9166666666666666
```

```
In [87]: precision_score(Y_test,y_pred1,average='macro')
```

```
Out[87]: 0.9629629629629629
```

```
In [88]: precision_score(Y_test,y_pred2,average='macro')
```

```
Out[88]: 0.9166666666666666
```

```
In [89]: recall_score(Y_test,y_pred2,average=None)
```

```
Out[89]: array([1.    , 0.875, 0.875])
```

```
In [90]: from sklearn.metrics import classification_report  
print (classification_report(Y_test, y_pred1))
```

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	14
versicolor	1.00	0.88	0.93	8
virginica	0.89	1.00	0.94	8
accuracy			0.97	30
macro avg	0.96	0.96	0.96	30
weighted avg	0.97	0.97	0.97	30

```
In [93]: from sklearn.metrics import classification_report  
print (classification_report(Y_test, y_pred2))
```

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	14
versicolor	0.88	0.88	0.88	8
virginica	0.88	0.88	0.88	8
accuracy			0.93	30
macro avg	0.92	0.92	0.92	30
weighted avg	0.93	0.93	0.93	30


```
In [91]: plt.xlabel('Features')
plt.ylabel('Species')


pltX = df.loc[:, 'sepal_length']
pltV = df.loc[:, 'species']
plt.scatter(pltX, pltV, color='red', label='sepal_length')

pltX = df.loc[:, 'sepal_width']
pltV = df.loc[:, 'species']
plt.scatter(pltX, pltV, color='blue', label='sepal_width')

pltX = df.loc[:, 'petal_length']
pltV = df.loc[:, 'species']
plt.scatter(pltX, pltV, color='yellow', label='petal_length')

pltX = df.loc[:, 'petal_width']
pltV = df.loc[:, 'species']
plt.scatter(pltX, pltV, color='green', label='petal_width')

plt.legend(loc=4, prop={'size':8})
plt.show()
```



```
In [10]: import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

# Creating a bar plot
sns.barplot(x=feature_imp, y=feature_imp.index)

# Add Labels to your graph
plt.xlabel('Feature Importance Score')
plt.ylabel('Features')
plt.title("Visualizing Important Features")
plt.legend()
plt.show()
```

No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

